

EN605.617 Module 6: Stretch problem

The file didn't state a `#include` statement for `stdio.h` which contains the `printf()` method.

The kernels perform multiplication and then adds result to the result destination. Which seems wrong unless the result matrix is zero. Mathematically it can be expressed as:

```
result[i] += multiplier[i]*multiplicand[i]
```

The shared memory portion seem to be more generic since `multiplicand` values are loaded into shared memory. The constant memory version seems to just declare a local array with pre-defined values from 0 to 15 and multiply the input by that.

I like how there are methods to transfer from global to shared memory. This is useful to see exactly how the transfer is being performed.

The code seems to be a bit disorganized, I had to jump around much.

A version of the code I modified is presented in the following pages. It compiles and runs on my local machine and it presents the same results I saw for the assignment program I made. Shared memory is overall faster. This is with both kernels using pinned memory.

```
#include <stdio.h>

// constant values that do not change at runtime
const unsigned int NUM_ROWS = 640*32;
const unsigned int NUM_COLS = 16;
const unsigned int NUM_SHARED_COEFF = 32;
const unsigned int NUM_REGISTER_COEFF = 16;
unsigned int num_threads = 256;

// from global_memory.cu file provided in Module 4 Vocareum lab
// create a timer object to test the duration of an operation
__host__ cudaEvent_t get_time(void)
{
    cudaEvent_t time;
    cudaEventCreate(&time);
    cudaEventRecord(time);
    return time;
}

// file previously declared matrix with default values. Each element contains the
// sum of the row and
// product index
void fill_matrix(unsigned int num_rows, unsigned int num_cols, unsigned int *
matrix_to_fill) {

    for (int col_index = 0; col_index < num_cols; col_index++)
    {
        for (int row_index = 0; row_index < num_rows; row_index++)
        {
            matrix_to_fill[row_index*(num_cols)+col_index] = row_index +
col_index;
        }
    }
}

// copy data from one type of device memory to another
__device__ void copy_btw_device_memory(unsigned int *mem_to_copy, unsigned int
*mem_to_fill, unsigned int num_elements)
{
    for (int elem_index = 0; elem_index < num_elements; elem_index++)
    {
        mem_to_fill[elem_index] = mem_to_copy[elem_index];
    }
}

// multiply matrices using shared memory (e.g. mx, where m is constant)
__device__ void multiply_values_shared(unsigned int * var_mat, unsigned int *
result_mat, unsigned int *shared_coeff_mat)
{

```

```

        unsigned int index = blockIdx.x * blockDim.x + threadIdx.x;
        atomicAdd(&result_mat[index / NUM_COLS], var_mat[index] *
shared_coeff_mat[index % NUM_COLS]);
    }

// store coefficients in shared memory based on thread location. Threads in the
top half of the
// coeff_mat use one set of coefficients, while threads in the bottom half use an
alternate set
__device__ void calc_shared_mem(unsigned int * coeff_mat, unsigned int *
shared_mat)
{
    int index = blockIdx.x * blockDim.x + threadIdx.x;
    unsigned int padding = (index / ((NUM_COLS*NUM_ROWS) / 2))*NUM_COLS;
    for (int i = 0; i < NUM_COLS; i++) {
        shared_mat[i] = coeff_mat[i + padding];
    }
}

// load data into shared memory and perform linear model multiplication (y=mx)
using shared memory
__global__ void linear_multiply_shared_mem_device(unsigned int * var_mat, unsigned
int * coeff_mat, unsigned int * result_mat)
{
    // declare shared memory, and load values into shared memory
    __shared__ unsigned int shared_coeff_mat[NUM_SHARED_COEFF];
    calc_shared_mem(coeff_mat, shared_coeff_mat);
    multiply_values_shared(var_mat, result_mat, shared_coeff_mat);
}

// perform linear model multiplication (y=mx) using constant memory
__global__ void linear_multiply_register_mem_device(unsigned int * var_mat,
unsigned int * result_mat)
{
    unsigned int index = blockIdx.x * blockDim.x + threadIdx.x;
    unsigned int register_mem[] = {0,1,2,3,4,5,6,7,8,9,10,11,12,13,14,15};
    atomicAdd(&result_mat[index / NUM_COLS], var_mat[index] * register_mem[index
% NUM_COLS]);
}

// print results of the GPU linear model multiplication
void print_select_results(unsigned int * data_to_print)
{
    printf("head of results \n");
    for (int i = 0; i < 10; i++) {
        printf("%i ", data_to_print[i]);
    }
    printf("\n transition between coefficients \n");
    for (int i = NUM_ROWS / 2 - 5; i < NUM_ROWS / 2 + 5; i++) {

```

```

        printf("%i ", data_to_print[i]);
    }
    printf(" tail of results \n");
    for (int i = NUM_ROWS - 10; i < NUM_ROWS; i++) {
        printf("%i ", data_to_print[i]);
    }
    printf("\n");
}

// perform QA/QC test for copying memory from GPU register memory to GPU global
memory
__global__ void test_register_copy(unsigned int * output_data, unsigned int
num_elements)
{
    unsigned int register_mem[] = {0,1,2,3,4,5,6,7,8,9,10,11,12,13,14,15};
    copy_btw_device_memory(register_mem, output_data, num_elements);
}

// perform QA/QC test for copying memory from global memory to GPU shared memory
back to GPU global memory
__global__ void test_shared_copy(unsigned int * input_data, unsigned int *
output_data, unsigned int num_elements)
{
    __shared__ unsigned int shared_int_data[16];
    copy_btw_device_memory(input_data, shared_int_data, num_elements);
    copy_btw_device_memory(shared_int_data, output_data, num_elements);
}

// test ability to copy to/from register memory
void test_register_mem_cpy() {
    unsigned int *h_results, *h_input, *d_results;
    cudaMalloc((void **)&d_results, sizeof(unsigned int)*16);
    cudaError_t status = cudaMallocHost((void **)&h_results, sizeof(unsigned int)
*16);
    if (status != cudaSuccess)
        printf("Error allocating pinned host memory\n"); // from Mark Harris,
Nvidia blogs
    status = cudaMallocHost((void **)&h_input, sizeof(unsigned int) * 16);
    if (status != cudaSuccess)
        printf("Error allocating pinned host memory\n"); // from Mark Harris,
Nvidia blogs

    fill_matrix(16, 1, h_input);

    test_register_copy<<< 1, 16 >>>(d_results, 16);

    cudaMemcpy(h_results, d_results, sizeof(unsigned int)*16,
cudaMemcpyDeviceToHost);

```

```

    printf("register memory copy test \n");
    for(int i = 0; i < 16; i++) {
        printf("expected: %i, actual: %i \n" , h_input[i], h_results[i]);
    }

    cudaFree(d_results);
    cudaFreeHost(h_input);
    cudaFreeHost(h_results);
}

// test ability to copy to/from shared memory
void test_shared_mem_cpy() {
    unsigned int *h_results, *h_input, *d_results, *d_input;
    cudaMalloc((void **)&d_results, sizeof(unsigned int)*16);
    cudaMalloc((void **)&d_input, sizeof(unsigned int)*16);
    cudaError_t status = cudaMallocHost((void**)&h_results, sizeof(unsigned int)
*16);
    if (status != cudaSuccess)
        printf("Error allocating pinned hot memory\n"); // from Mark Harris,
NVidia blogs
    status = cudaMallocHost((void**)&h_input, sizeof(unsigned int) * 16);
    if (status != cudaSuccess)
        printf("Error allocating pinned hot memory\n"); // from Mark Harris,
NVidia blogs

    fill_matrix(16, 1, h_input);

    cudaMemcpy(d_input, h_input, sizeof(unsigned int)*16,
cudaMemcpyHostToDevice);
    test_shared_copy<<< 1, 16 >>>(d_input,d_results,16);

    cudaMemcpy(h_results, d_results, sizeof(unsigned int)*16,
cudaMemcpyDeviceToHost);

    printf("shared memory copy test \n");
    for(int i = 0; i < 16; i++) {
        printf("expected: %i, actual: %i \n" , h_input[i], h_results[i]);
    }

    cudaFree(d_results);
    cudaFree(d_input);
    cudaFreeHost(h_input);
    cudaFreeHost(h_results);
}

// host and device operations to perform linear multiplication (e.g. mx) on the
device using register memory

```

```

float linear_multiply_register_mem_host(bool debug, int num_thrd)
{
    num_threads = num_thrd;

    // declare and allocate device memory
    unsigned int *d_var;
    unsigned int *d_results;
    cudaMalloc((void **)&d_var, sizeof(unsigned int)* (NUM_COLS*NUM_ROWS));
    cudaMalloc((void **)&d_results, sizeof(unsigned int)*NUM_ROWS);

    // declare and allocate pinned host memory
    unsigned int *h_var, *h_results, *h_coeff;
    cudaError_t status = cudaMallocHost((void **)&h_var, sizeof(unsigned int)
*(NUM_ROWS*NUM_COLS));
    if (status != cudaSuccess)
        printf("Error allocating pinned host memory\n"); // from Mark Harris,
Nvidia blogs
    status = cudaMallocHost((void **)&h_results, sizeof(unsigned int) *NUM_ROWS);
    if (status != cudaSuccess)
        printf("Error allocating pinned host memory\n"); // from Mark Harris,
Nvidia blogs
    status = cudaMallocHost((void **)&h_coeff, sizeof(unsigned int)
*(NUM_REGISTER_COEFF));
    if (status != cudaSuccess)
        printf("Error allocating pinned host memory\n"); // from Mark Harris,
Nvidia blogs

    // fill declared matrices with default values (row + column index)
    fill_matrix(NUM_ROWS, NUM_COLS, h_var);
    fill_matrix(NUM_REGISTER_COEFF, 1, h_coeff);

    // measure how long it takes to copy data to device, run a kernel, and copy the
data back to host
    cudaEvent_t start_time = get_time();

    // copy data to device
    cudaMemcpy(d_var, h_var, sizeof(unsigned int)* (NUM_COLS*NUM_ROWS),
cudaMemcpyHostToDevice);
    cudaMemcpy(d_results, h_results, sizeof(unsigned int)*(NUM_ROWS),
cudaMemcpyHostToDevice);

    // perform linear multiplication on GPU and copy data from GPU to host
    linear_multiply_register_mem_device <<< (NUM_ROWS*NUM_COLS + num_threads - 1)
/ num_threads, num_threads >>> (d_var, d_results);
    cudaMemcpy(h_results, d_results, sizeof(unsigned int)*NUM_ROWS,
cudaMemcpyDeviceToHost);

    // finish timing performance and record result

```

```

    cudaEvent_t end_time = get_time();
    cudaEventSynchronize(end_time);
    float delta = 0;
    cudaEventElapsedTime(&delta, start_time, end_time);

    // print select rows of results for debugging, quality assurance purposes
    if (debug) {
        print_select_results(h_results);
    }

    // free device memory
    cudaFree(d_var);
    cudaFree(d_results);

    // free pinned memory
    cudaFreeHost(h_var);
    cudaFreeHost(h_results);
    cudaFreeHost(h_coeff);

    return delta;
}

// host and device operations to perform linear multiplication (e.g. mx) on the
// device using shared memory
float linear_multiply_shared_mem_host(bool debug, int num_thrd)
{
    num_threads = num_thrd;

    // declare and allocate device memory
    unsigned int * d_coeff, *d_var, *d_results;
    cudaMalloc((void **)&d_coeff, sizeof(unsigned int) * NUM_SHARED_COEFF);
    cudaMalloc((void **)&d_var, sizeof(unsigned int) * (NUM_COLS*NUM_ROWS));
    cudaMalloc((void **)&d_results, sizeof(unsigned int)*NUM_ROWS);

    // declare and allocate pinned host memory
    unsigned int *h_var, *h_results, *h_coeff;
    cudaError_t status = cudaMallocHost((void**)&h_var, sizeof(unsigned int)
*(NUM_ROWS*NUM_COLS));
    if (status != cudaSuccess) printf("Error allocating pinned hot memory\n"); //
from Mark Harris, NVidia blogs
    status = cudaMallocHost((void**)&h_results, sizeof(unsigned int) *NUM_ROWS);
    if (status != cudaSuccess) printf("Error allocating pinned hot memory\n"); //
from Mark Harris, NVidia blogs
    status = cudaMallocHost((void**)&h_coeff, sizeof(unsigned int)
*NUM_REGISTER_COEFF);
    if (status != cudaSuccess) printf("Error allocating pinned hot memory\n"); //
from Mark Harris, NVidia blogs

```

```

    // fill declared matrices with default values (row + column index)
    fill_matrix(NUM_ROWS, NUM_COLS, h_var);
    fill_matrix(NUM_SHARED_COEFF, 1, h_coeff);

    // measure how long it takes to copy data to device, run a kernel, and copy
the data back to host
    cudaEvent_t start_time = get_time();

    // copy data to device, including constant memory
    cudaMemcpy(d_coeff, h_coeff, sizeof(unsigned int) * NUM_SHARED_COEFF,
cudaMemcpyHostToDevice);
    cudaMemcpy(d_var, h_var, sizeof(unsigned int)* (NUM_COLS*NUM_ROWS),
cudaMemcpyHostToDevice);
    cudaMemcpy(d_results, h_results, sizeof(unsigned int)*(NUM_ROWS),
cudaMemcpyHostToDevice);

    // perform linear multiplication on GPU and copy data from GPU to host
    linear_multiply_shared_mem_device <<< (NUM_ROWS*NUM_COLS + num_threads - 1) /
num_threads, num_threads >>> (d_var, d_coeff, d_results);
    cudaMemcpy(h_results, d_results, sizeof(unsigned int)*NUM_ROWS,
cudaMemcpyDeviceToHost);

    // finish timing performance and record result
    cudaEvent_t end_time = get_time();
    cudaEventSynchronize(end_time);
    float delta = 0;
    cudaEventElapsedTime(&delta, start_time, end_time);

    // print select rows of results for debugging, quality assurance purposes
    if (debug) { print_select_results(h_results);}

    // free device memory
    cudaFree(d_coeff);
    cudaFree(d_var);
    cudaFree(d_results);

    // free pinned memory
    cudaFreeHost(h_var);
    cudaFreeHost(h_results);
    cudaFreeHost(h_coeff);

    return delta;
}

// process input args to determine if debug statements should be printed
bool set_debug_flag(int argc, char** argv) {
    bool debug = false;
    if (argc>1)
    {

```



```
        if (argv[1][0] == 't' or argv[1][0] == 'T')
        {
            debug = true;
            printf("run debug \n");
        }
    }
    return(debug);
}

int main(int argc, char** argv)
{
    float duration;

    // process input and determine if debug statements should be printed
    bool should_debug = set_debug_flag(argc, argv);

    test_register_mem_cpy();
    test_shared_mem_cpy();

    // perform linear multipliation (y =mx) using shared memory
    duration = linear_multiply_shared_mem_host(should_debug, 256);
    printf("time from copy to completion for linear multiply algorithm with
shared memory: %f \n", duration);

    // perform linear multipliation (y=mx) using register memory
    duration = linear_multiply_register_mem_host(should_debug, 256);
    printf("time from copy to completion for linear multiply algorithm with
register memory: %f \n", duration);

    printf("finished program");
    return 1;
}
```

Results:

register memory copy test

expected: 0, actual: 0

expected: 1, actual: 1

expected: 2, actual: 2

expected: 3, actual: 3

expected: 4, actual: 4

expected: 5, actual: 5

expected: 6, actual: 6

expected: 7, actual: 7

expected: 8, actual: 8

expected: 9, actual: 9

expected: 10, actual: 10

expected: 11, actual: 11

expected: 12, actual: 12

expected: 13, actual: 13

expected: 14, actual: 14

expected: 15, actual: 15

shared memory copy test

expected: 0, actual: 0

expected: 1, actual: 1

expected: 2, actual: 2

expected: 3, actual: 3

expected: 4, actual: 4

expected: 5, actual: 5

expected: 6, actual: 6

expected: 7, actual: 7

expected: 8, actual: 8

expected: 9, actual: 9

expected: 10, actual: 10

expected: 11, actual: 11

expected: 12, actual: 12

expected: 13, actual: 13

expected: 14, actual: 14

expected: 15, actual: 15

time from copy to completion for linear multiply algorithm with shared memory:

1.332064

time from copy to completion for linear multiply algorithm with register memory:

1.599648

finished program